



Pipeline Component Breakdown
(single layer, 64 MiB chunks)



H2D PCIe Interference
(lower = better)



Full Model (Qwen3-8B, RTX 4090)
end-to-end training

